

DATA ANALYTICS AND VISUALIZATION

(24DSB3101)

Machine Learning-Based Healthcare Data Analytics for Early Disease Risk Prediction

1. Problem Statement

Healthcare datasets contain demographic, clinical, and lifestyle attributes that can help identify individuals at elevated risk of disease. However, missing values, inconsistent records, class imbalance, and complex relationships among health factors can make reliable prediction challenging. Recent research also identifies concerns around model interpretability, dataset diversity, clinical workflow integration, and real-world validation. This project proposes an end-to-end healthcare analytics framework that preprocesses benchmark health datasets, develops disease-risk prediction models, evaluates them using standard metrics, and presents risk assessment through an interactive dashboard.

2. Literature Survey

Recent and foundational papers were selected to understand predictive healthcare, explainable AI, disease-risk modelling, interpretability, and clinical decision-support research. The survey includes foundational XAI literature from 2016–2020 followed by recent 2024–2025 healthcare studies. The recent studies are ordered with 2024 first, followed by 2025.

No.	Research Paper	Year	Main Focus	Limitations
1	<i>"Why Should I Trust You?": Explaining the Predictions of Any Classifier</i>	2016	Introduces LIME, a model-agnostic technique for generating local, interpretable explanations for black-box classifier predictions.	Focuses mainly on local explanations and is not a healthcare-specific, multi-disease prediction framework.
2	<i>A Unified Approach to Interpreting Model Predictions</i>	2017	Introduces SHAP, a unified framework that assigns feature-importance values to individual predictions and supports interpretation of complex models.	Primarily addresses model interpretation; computational cost and complete clinical decision-support integration remain challenges.
3	<i>Towards A Rigorous Science of Interpretable Machine Learning</i>	2017	Defines interpretability concepts and proposes a taxonomy and directions for rigorous evaluation of interpretable machine-learning systems.	Primarily methodological and conceptual; it does not provide an end-to-end healthcare prediction and explanation workflow.
4	<i>Causability and Explainability of Artificial Intelligence in Medicine</i>	2019	Discusses explainability and causability in medical AI, emphasizing transparency, traceability, and the quality of explanations for human users.	Conceptual focus with limited implementation of an integrated multi-disease prediction, XAI, and interactive visualization system.
5	<i>Explainability for Artificial Intelligence in Healthcare: A Multidisciplinary Perspective</i>	2020	Examines explainability in healthcare from technological, medical, legal, ethical, and patient perspectives, especially for clinical decision support.	Highlights broad adoption challenges but does not provide a reusable multi-disease prediction framework with integrated SHAP/LIME visualization.
6	<i>How Explainable Artificial Intelligence Can Increase or Decrease Clinicians' Trust in AI Applications in Health Care: Systematic Review</i>	2024	Systematically reviews empirical evidence on how XAI explanations affect clinicians' trust in AI-driven clinical decision-making.	Explanation quality can increase or decrease trust; excessive trust, unclear explanations, and limited evidence on optimal clinical integration remain concerns.
7	<i>The Role of Explainable Artificial Intelligence in Disease Prediction: A Systematic Literature Review and Future Research Directions</i>	2025	Reviews XAI techniques used for disease prediction and identifies major research challenges related to explainability and disease-risk modelling.	Limited dataset diversity, model complexity, and dependence on single data types restrict generalizability.
8	<i>The Application of Explainable Artificial Intelligence in the Prediction, Diagnoses, Treatment, and Management of Chronic Diseases: A Systematic Review</i>	2025	Reviews XAI applications across chronic diseases and highlights the use of methods such as SHAP and LIME.	Limited multimodal data, limited data volume, and insufficient real-world clinical validation.
9	<i>Presenting Artificial Intelligence Predictions Based on Electronic Medical Records to Clinicians in Hospitals: A Systematic Review</i>	2025	Reviews how AI predictions derived from structured electronic medical records are presented to clinicians in hospital settings.	Limited evidence on optimal presentation methods and practical clinical integration.
10	<i>Towards an Evaluation Framework for Explainable Artificial Intelligence Systems for Health and Well-being</i>	2025	Proposes an evaluation framework for explainable AI systems in health and well-being, focusing on systematic assessment of XAI quality.	Evaluation remains challenging, particularly for transparency, explanation quality, and human understanding of AI decisions.

3. Research Gap Identification

- **Gap 1 — Disease-specific approaches:** Many existing studies concentrate on one disease or one dataset, making it difficult to establish a reusable healthcare analytics workflow.
- **Gap 2 — Limited data diversity:** Recent literature identifies limited dataset diversity and reliance on single data types as important limitations in disease-prediction research.
- **Gap 3 — Prediction vs. interpretability:** Many systems prioritize predictive performance, while explainability, transparency and trust are treated as secondary components.
- **Gap 4 — Lack of integrated XAI workflow:** SHAP/LIME are frequently applied individually, but there is scope for a standardized workflow combining prediction, global explanation, local explanation and visualization.
- **Gap 5 — Limited practical presentation:** Research often reports model performance without providing an accessible interface through which users can inspect risk predictions and their contributing factors.

4. Objectives

- To analyze healthcare datasets using Pandas, NumPy, Matplotlib, Seaborn, and automated profiling techniques to understand data distributions, correlations, data-quality issues, and significant disease-risk factors.
- To develop disease-risk prediction models using Logistic Regression, Decision Tree, Random Forest and XGBoost, and evaluate them using Accuracy, Precision, Recall, F1-Score and ROC-AUC.
- To improve the interpretability of disease-risk predictions using SHAP-based global and local feature explanations, enabling identification of the major factors influencing individual predictions.
- To develop an interactive Streamlit-based healthcare analytics dashboard that presents disease-risk probabilities, important risk factors, model performance and explainable prediction results for decision support.
- To evaluate the proposed healthcare data analytics framework comprehensively based on data quality, prediction performance, model interpretability, and overall system effectiveness.

5. Innovation

- End-to-End Healthcare Analytics – Integrates data preprocessing, analysis, prediction, and visualization.
- Multi-Model Prediction – Compares multiple ML models with XGBoost for disease-risk prediction.
- Explainable AI – Uses SHAP to identify the factors influencing each prediction.
- Interactive Dashboard – Uses Streamlit to provide an easy-to-use healthcare prediction interface.
- Visual Risk Insights – Presents risk probability, important features, and analytical charts clearly.

6. Novelty

- Unified Healthcare Analytics Framework – Integrates data preprocessing, EDA, feature engineering, ML prediction, explainability, and visualization in one workflow.
- Multi-Disease Risk Prediction – Applies the framework to multiple healthcare datasets such as heart disease, diabetes, and stroke.
- Explainable Risk Prediction – Uses SHAP to identify the major health factors contributing to each prediction.
- Risk-to-Insight Analysis – Provides both the disease-risk probability and the key factors influencing the predicted risk.

- **Interactive Decision-Support Dashboard** – Uses Streamlit to present predictions, visual analytics, model performance, and explanations in one interface.
- **Systematic Model Evaluation** – Evaluates multiple ML models using Accuracy, Precision, Recall, F1-Score, and ROC-AUC to select the most suitable model.

7. Feasibility

Technical Feasibility

The project can be implemented using publicly available healthcare datasets and Python-based open-source libraries such as Pandas, Scikit-learn, XGBoost, SHAP and Streamlit.

Data Feasibility

Public benchmark datasets are available for heart disease, diabetes and stroke prediction, making initial experimentation possible without requiring private hospital records.

Implementation Feasibility

The system can initially be developed using Jupyter Notebook/VS Code and subsequently integrated into a Streamlit dashboard.

Project Phases

Phase 1 → Literature Survey & Dataset Collection

Phase 2 → Data Cleaning & EDA

Phase 3 → Feature Engineering

Phase 4 → ML Model Development

Phase 5 → Model Evaluation

Phase 6 → SHAP Explainability

Phase 7 → Streamlit Dashboard

Phase 8 → Testing, Documentation & Final Report

8. Project Plan

Phase	Phase Name	Description
P1	Literature Survey & Problem Analysis	Study recent healthcare ML and XAI research papers, understand existing approaches, identify limitations, and finalize the research gap and problem statement.
P2	Dataset Collection & Data Understanding	Collect benchmark healthcare datasets such as UCI Heart Disease, Diabetes, and Stroke datasets and analyze their attributes, target variables, data types, and quality.
P3	Data Preprocessing & Exploratory Data Analysis	Handle missing values, duplicates and outliers, perform encoding/scaling where required, and analyze distributions, correlations and important health-risk patterns using EDA techniques.
P4	Feature Engineering & ML Model Development	Select or create relevant features and develop disease-risk prediction models using Scikit-learn algorithms and XGBoost.
P5	Model Evaluation & Explainability	Evaluate models using Accuracy, Precision, Recall, F1-Score and ROC-AUC, select suitable models, and use SHAP to explain important factors behind predictions.
P6	Dashboard Development & System Integration	Integrate the selected ML model and SHAP explanations into a Streamlit dashboard to display disease-risk probability.

9. Architecture

Benchmark Healthcare Datasets

Collect healthcare datasets such as UCI Heart Disease, Diabetes, and Stroke datasets for model development.

Data Preprocessing

Clean the collected data by handling missing values, duplicates, outliers, categorical values, and scaling where required.

Exploratory Data Analysis (EDA)

Analyze the datasets using statistics, distributions, correlations, and visualizations to identify important health-risk patterns.

Feature Engineering & Selection

Create or transform useful features and select the most relevant health attributes for prediction.

Train-Test Split

Divide the processed dataset into training and testing data to train and evaluate the models fairly.

Machine Learning Models

Develop prediction models using Logistic Regression, Decision Tree, Random Forest, and XGBoost.

Performance Evaluation

Compare the models using Accuracy, Precision, Recall, F1-Score, ROC-AUC, and Confusion Matrix.

Explainable AI (XAI)

Use SHAP to explain which health factors contribute most to the model's risk prediction.

Visualization & Dashboard

Use Matplotlib, Seaborn, and Streamlit to present predictions, risk factors, model results, and visual insights.

Final Output / Decision Support

Provide an early disease-risk prediction, probability score, and important contributing factors through the dashboard.

Team Members (Batch – 7)

E. Supraja - 2420030357

M. Varshitha - 2420030479

C. Keerthana - 2420030581

CH. Bhargavi - 2420030588